

Know the bill before you sign for it

Governed doesn't have to mean expensive — but it does have to mean *knowable*. The same agent architecture runs from roughly a hundred dollars a month in a disposable sandbox to several thousand in a zone-redundant production estate, and the difference is four or five variables rather than a redesign.

Sandbox demo ~\$100 per month · public endpoints, Consumption APIM, Basic search. Non-sensitive demo data only.	Module defaults ~\$680 per month · private endpoints on, Developer APIM, Standard search, 4 governed agents.	Production ~\$3.6k per month idle · Premium APIM in Internal VNet, S1 search x3, policies enforced.
---	--	---

01 Where the money goes

At the module defaults — private endpoints on, four agents, ~20K conversations a month on gpt-4o-mini. Two lines dominate; the rest is rounding.

AI Search	\$490
Private endpoints	\$57
API Management	\$50
Log Analytics	\$21
Defender for Cloud	\$21
OpenAI tokens	\$16
Content Safety	\$15
Storage + Key Vault	\$5
Networking	\$5
\$0	
\$490 / month	

92% of that bill arrives whether or not anyone logs in. AI Search and API Management bill while they *exist*, idle or not. Token spend is the opposite — \$0 until an agent is actually used, because deployment quota is a rate limit, not a reservation.

02 The five knobs that move the needle

- API Management SKU.** Developer ~\$50 (no SLA) versus Premium ~\$2,800 per unit. Premium buys Internal VNet *and* availability zones — if you only need VNet, Standard v2 is ~\$700.
- AI Search tier x units.** Basic 1x1 ~\$75 versus Standard 1x3 ~\$735. Replicas buy query SLA; partitions buy index size. Scale the one you actually need.
- Private endpoints.** ~\$9.50/mo each, so ~\$57 for a typical six. Turning them off saves that and costs you network isolation — demo data only.
- Log Analytics retention and ingest.** 5 GB/mo free, ~\$2.53/GB after, plus ~\$0.12/GB/mo past 31 days. Reuse the platform workspace rather than standing up your own.
- Environments and agents.** Mostly licensing and tenant capacity rather than Azure spend — but each production Dataverse environment draws ~1 GB from tenant capacity, with overage around \$40/GB/mo.

Cheapest honest evaluation: the sandbox profile — ~\$100 for a month, or a few dollars torn down the same afternoon. It runs the full twelve-proof test suite for about five cents in tokens.

03 Profile comparison

SETTING	SANDBOX	DEFAULTS	PROD
Private endpoints	off	on	on
APIM SKU	Consum.	Dev	Premium
AI Search	basic 1x1	std 1x2	std 1x3
PP environments	1 sandbox	3	3
Governed agents	1	4	4
LAW retention	30 d	90 d	90 d
Defender AI plan	off	on	on
Policies	Audit	Enforced	Enforced
Idle / month	~\$100	~\$680	~\$3.6k

04 Metered, not fixed

SERVICE	RATE
gpt-4o	\$2.50 / \$10 per 1M
gpt-4o-mini	\$0.15 / \$0.60 per 1M
text-embedding-3-large	\$0.13 per 1M
Content Safety	~\$0.75 / 1K records
APIM Consumption	~\$3.50 / 1M calls

05 Billed outside Azure

Power Platform and Copilot Studio are licensing and tenant capacity, not Azure metering — they never appear on the Azure invoice. Premium Power Platform runs ~\$5–20 per user per month; Copilot Studio runs on a licence or message packs at roughly \$200 per 25,000 messages.

You can turn it off. terraform destroy per module, in reverse order. Destroy AI Search and API Management first — those are what stop the meter. Key Vault purge protection holds a soft-deleted vault, and its name, for 90 days with no force-purge, which is exactly why the sandbox profile sets it to `false` : apply, demo, destroy, rebuild the same day.

Methodology. Figures are list prices in USD for East US 2, pay-as-you-go, with no EA or CSP discount, at early-2026 rates, and mirror `azure/COSTS.md` in the Adapt Cloud reference implementation. "Idle" means the cost incurred with the platform deployed and nobody using it. Estimates only — confirm against the Azure Pricing Calculator before committing budget. · **Adapt Cloud** — governed AI agent platforms on Azure.