

Your agents are identities. Govern them like it.

Most agent projects die in the gap between “it answered the question” and “it can touch production data.” Adapt Cloud closes that gap with infrastructure rather than policy documents: agents deploy inside a real Azure landing zone, each with its own Entra identity, behind guardrails that *deny* the bad deployment instead of reporting on it next quarter.

Agents as first-class identities

One Entra Agent ID or user-assigned managed identity per agent. No shared service principals, no impersonated users, no client secrets — federated credentials instead.

A real landing zone

An Application landing zone under an Application Platform management group — created idempotently if your tenant doesn’t already have one, then peered to your connectivity hub.

Policy as code, in Deny

Seven definitions in one initiative, assigned at management-group scope — so it covers the workloads nobody told you about.

Low-code first

Power Platform Managed Environments with DLP classifying the full connector catalog into Business, Non-Business and Blocked, plus APIM-fronted custom connectors from your own OpenAPI specs.

Traceable by design

Central Log Analytics, per-agent Application Insights, diagnostic settings enforced by policy, Defender for Cloud AI threat protection, Purview/DSPM hooks.

Human in the loop

Sensitive operations sit behind an APIM policy requiring an approver header. An agent can prepare a payout. It cannot issue one alone.

01 Twelve claims, twelve tests

Architecture diagrams are easy to draw and impossible to falsify. Every pillar here is paired with a runnable test designed to fail if the control is removed.

THE CLAIM	PROVEN BY
✓ Per-agent identity, no shared keys	unauth call → 401/403
✓ APIM validates the Entra JWT	federated token → 200
✓ Tenant-wide rate limit holds	601-call burst → 429s
✓ Connectors match the OpenAPI contract	FNOL → triage → payout
✓ Approved models only	gpt-3.5-turbo → denied
✓ No public network access on AI/PaaS	public endpoint → denied
✓ Required governance tags	missing owner → denied
✓ Prompt Shields catch jailbreaks	10 adversarial prompts refused
✓ Domain accuracy of the agent	10 labelled FNOL cases scored
✓ Human-in-the-loop on payouts	no approver header → 403
✓ Calls traceable to the agent identity	App Insights trace by agent name
✓ Initiative enforced at MG scope	RequestDisallowedByPolicy

The end-to-end script chains all of it — apply the sandbox profile, wait for the first policy scan, run the contract tests, run the evaluations, prove the denials, print a compliance summary from `az policy state list`, then destroy. About five cents in tokens.

02 What gets deployed

- **Application Platform landing zone** — management group, vended workload subscription, spoke network, central Log Analytics, platform managed identities.
- **Policy-as-code initiative** — allowed regions, approved model deployments, deny public network access, deny local auth on Cognitive Services, require diagnostic settings, require ownership and classification tags, audit managed identity.
- **Insurance workload landing zone** — Power Platform Managed Environment and connector DLP, APIM-fronted API connectors, Copilot Studio / AI Foundry agents each with their own identity, Key Vault, Application Insights, private endpoints.
- **Control-mapping matrix** — control-by-control traceability from each CAF recommendation to the file implementing it.

03 What it costs

A disposable sandbox is roughly **\$100/month**; the module defaults land near **\$680**; a zone-redundant production posture starts around **\$3.6k/month** idle. AI Search and API Management dominate — and both bill while they exist, idle or not. Token spend is \$0 until an agent is used. Full model in the FinOps one-pager.

04 How we work

1 WEEK

Assess

Which identities your agents run under, what they can reach, and a gap map against the CAF governance disciplines.

3-4 WEEKS

Land

Landing zone, initiative in Audit, one real workload agent deployed end to end. You keep the Terraform.

ONGOING

Enforce

Audit → Deny once the estate is clean, tests wired into your pipeline, operating model handed over.

Start with the sandbox. We stand up the dev-demo profile in your tenant, run the twelve proofs, walk your security and platform teams through what passed and what didn’t — and tear it down the same day if you want.

Adapt Cloud — governed AI agent platforms on Azure. Reference implementation: Application Platform landing zone, policy-as-code guardrails, per-agent Entra identities, Power Platform environments governed by connector-classification DLP, and an end-to-end test suite pairing every architectural claim with a runnable proof. Nothing here is meant to go straight to production — review, parameterise, and run it through your own pipeline first.